

✓ Importing Libraries and Data

```
import pandas as pd
import numpy as np

# Exploratory Data Analysis (EDA)
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt

# Modelling
import scipy.cluster.hierarchy as sch
from sklearn.cluster import KMeans
from sklearn.cluster import DBSCAN
from sklearn.cluster import AgglomerativeClustering
from sklearn.metrics import silhouette_score
from sklearn.metrics import davies_bouldin_score

df_segmentation = pd.read_csv("segmentation data.csv")
```

Project Objectives

The primary objective of this project is to perform a customer segmentation analysis using machine learning techniques. By clustering customers based on their demographic and behavioral characteristics, we aim to uncover distinct groups within the customer base. This segmentation will enable us to:

- **Understand Customer Profiles:** Identify and understand the different types of customers, including their purchasing behaviors and preferences.
- **Personalize Marketing Strategies:** Develop targeted marketing campaigns and personalized strategies for each customer segment to improve customer engagement and satisfaction.
- **Enhance Business Decision-Making:** Provide valuable insights for data-driven decision-making in areas such as product development, pricing strategies, and customer service enhancements.
- **Improve Resource Allocation:** Allocate resources more effectively by focusing on high-value customer segments and optimizing marketing and operational efforts.
- **Increase Revenue and Profitability:** Ultimately, drive business growth by increasing customer retention, enhancing customer lifetime value, and acquiring new customers through more effective segmentation strategies.

By achieving these objectives, the project will contribute to a deeper understanding of our customer base, leading to more informed business strategies and improved overall performance.

✓ Exploratory Data Analysis (EDA)

- **Data Understanding:** check the data types of each column to ensure they are appropriate for the analysis and calculate basic statistics (mean, median, standard deviation, etc.) to get an overview of the data distribution.
- **Data Cleaning:** identify and handle missing values through imputation or removal, depending on the context, and detect and address outliers that might skew the analysis.
- **Data Transformation:** convert categorical variables (Gender) into numerical format.
- **Data Visualization:** use scatter plots, and histograms to explore relationships between pairs of variables.

```
df_segmentation.head()
```



	CustomerID	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	1	Male	19	15	39
1	2	Male	21	15	81
2	3	Female	20	16	6
3	4	Female	23	16	77
4	5	Female	31	17	40

```
df_segmentation.dtypes
```

```

CustomerID      int64
Gender          object
Age            int64
Annual Income (k$)  int64
Spending Score (1-100)  int64
dtype: object

```

```
df_segmentation.describe()
```

```

CustomerID      Age  Annual Income (k$)  Spending Score (1-100)
count  200.000000  200.000000          200.000000          200.000000
mean    100.500000   38.850000           60.560000           50.200000
std     57.879185   13.969007           26.264721           25.823522
min      1.000000   18.000000           15.000000            1.000000
25%     50.750000   28.750000           41.500000           34.750000
50%     100.500000  36.000000           61.500000           50.000000
75%     150.250000  49.000000           78.000000           73.000000
max     200.000000  70.000000          137.000000           99.000000

```

```
# Searching for null values
```

```

null_counts = df_segmentation.isnull().sum()
total_rows = len(df_segmentation)
null_percentage = (null_counts / total_rows) * 100

```

```

null_summary_df = pd.DataFrame({
    'Column': null_counts.index,
    'Total Values': total_rows,
    'Null Values': null_counts.values,
    'Percentage': null_percentage.values
})
print(null_summary_df)

```

```

Column  Total Values  Null Values  Percentage
0      CustomerID      200          0         0.0
1         Gender      200          0         0.0
2          Age       200          0         0.0
3  Annual Income (k$)   200          0         0.0
4  Spending Score (1-100)  200          0         0.0

```

```
# Format Gender to binary
```

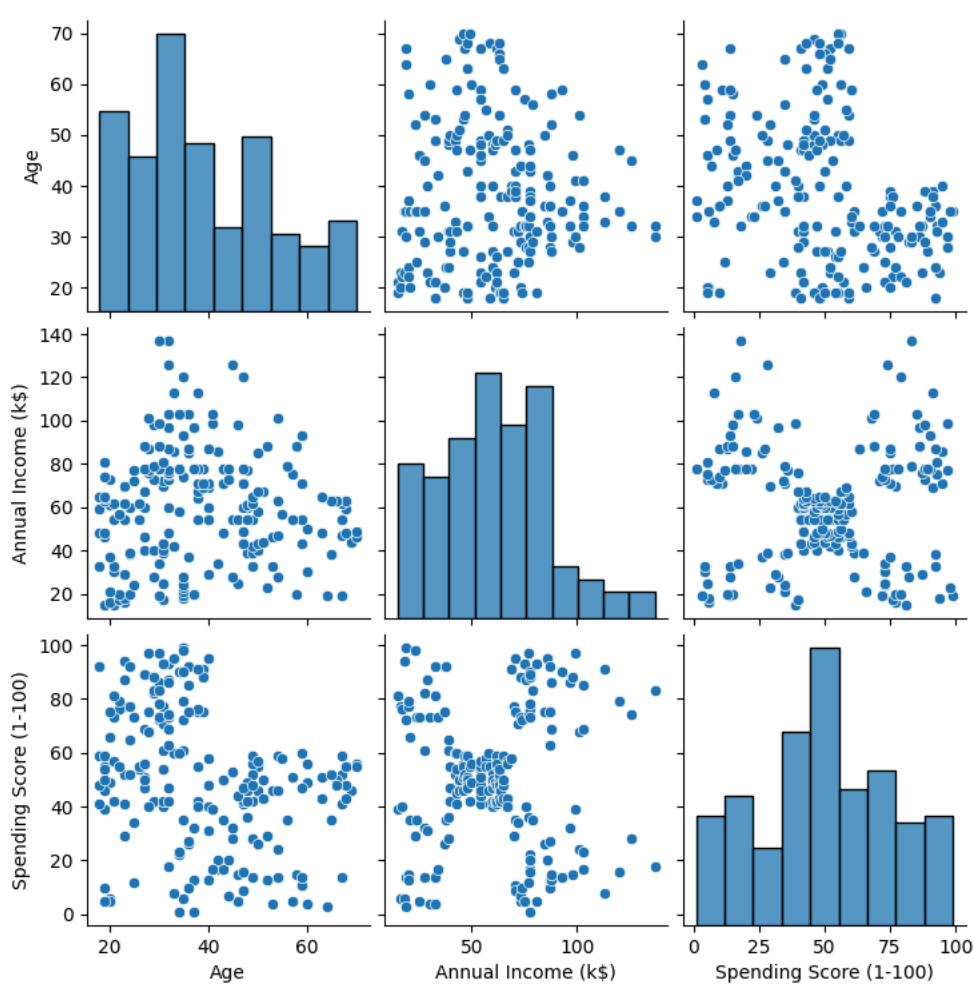
```
df_segmentation['Gender'] = df_segmentation['Gender'].map({'Male': 0, 'Female': 1})
```

```
# Distribution of numerical variables
```

```


sns.pairplot(df_segmentation[['Age', 'Annual Income (k$)', 'Spending Score (1-100)']])
plt.show()

```



```
# Gender distribution ==> Male = 0 and Female = 1
```

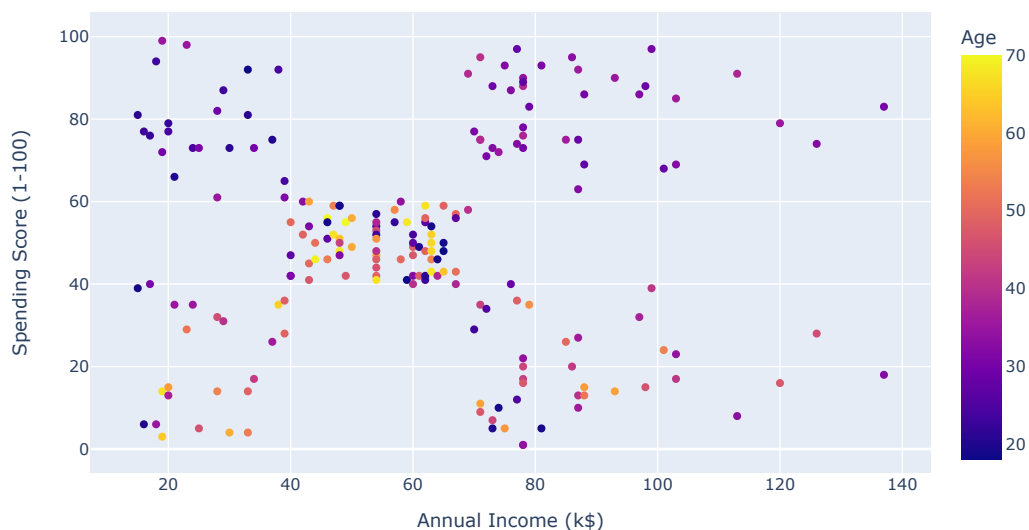
```
sns.countplot(x='Gender', data=df_segmentation)
plt.show()
```



```
fig = px.scatter(df_segmentation, x = 'Annual Income (k$)', y = 'Spending Score (1-100)', width = 800, color = "Age", title = 'Income v:
fig.show()
```




Income vs. Score by Age



```
fig = px.scatter(df_segmentation, x = 'Annual Income (k$)', y = 'Age', width = 800, color = "Gender",title = 'Income vs. Age by Gender')
fig.show()
```



Income vs. Age by Gender



Modelling


- **Model Selection:** choose appropriate clustering algorithms based on the data characteristics and analysis goals (K-means, DBSCAN, Agglomerative Clustering) experimenting with different parameters (number of clusters in K-means and AC, eps in DBSCAN) to find the optimal configuration.
- **Model Training:** apply the algorithm to the data to create the model.
- **Model Evaluation:** evaluate the quality of clusters using metrics like silhouette score and Davies-Bouldin index. After visualize the resulting clusters using scatter plots to understand their structure and distribution.

K-Means

```
# Determine inertia for different numbers of clusters

inertia = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init = 10)
    kmeans.fit(df_segmentation)
    inertia.append(kmeans.inertia_)

# Plot the elbow graph
plt.figure(figsize=(8, 5))
plt.plot(range(1, 11), inertia, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.title('Elbow Method')
plt.show()
```



```
# Setting K-Means

features_kmeans = df_segmentation.iloc[:, 1:] # Remove Id
k = 5
kmeans = KMeans(n_clusters=k, init='k-means++', random_state=42,n_init = 10)
Clusters_KMeans = kmeans.fit_predict(features_kmeans)
features_kmeans.loc[:, 'Cluster'] = Clusters_KMeans
```

```
# Analyze cluster characteristics
```

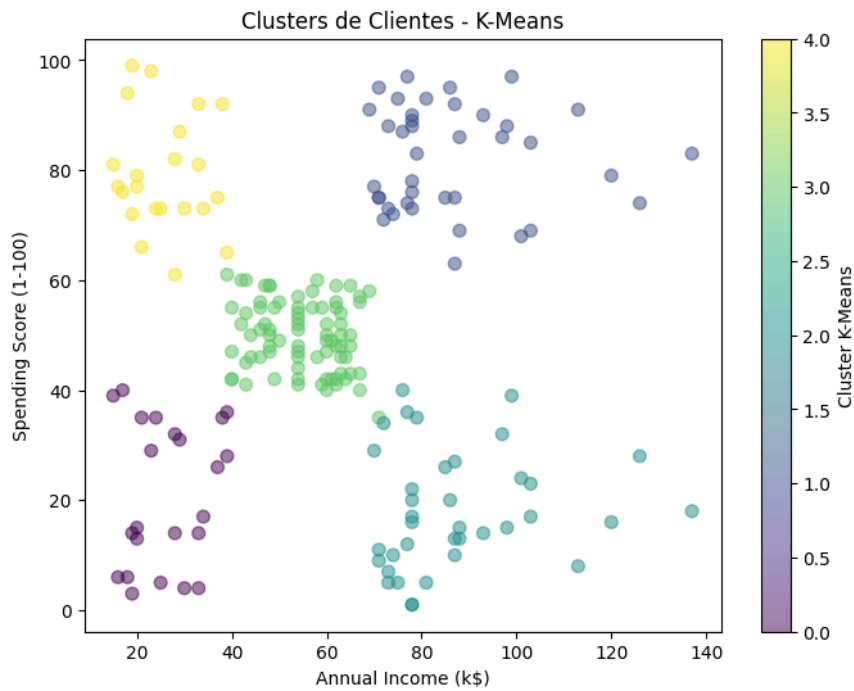
```
cluster_analysis = features_kmeans.groupby('Cluster').mean()
print(cluster_analysis)
```



Cluster	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	0.608696	45.217391	26.304348	20.913043
1	0.538462	32.692308	86.538462	82.128205
2	0.486486	40.324324	87.432432	18.189189
3	0.582278	43.126582	54.822785	49.835443
4	0.590909	25.272727	25.727273	79.363636

```
# Cluster Plot
```

```
plt.figure(figsize=(8, 6))
plt.scatter('Annual Income (k$)', 'Spending Score (1-100)', data=features_kmeans, c='Cluster', cmap='viridis', s=50, alpha=0.5)
plt.title('Clusters de Clientes - K-Means')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.colorbar(label='Cluster K-Means')
plt.show()
```



DBSCAN

```
# Setting DBSCAN
```

```
features_DBSCAN = df_segmentation.iloc[:, 1:] # Remove Id
dbscan = DBSCAN(eps=13, min_samples=10)
Clusters_DBSCAN = dbscan.fit_predict(features_DBSCAN)
features_DBSCAN.loc[:, 'Cluster'] = Clusters_DBSCAN
```

```
# Analyze cluster characteristics
```


```
cluster_analysis = features_DBSCAN.groupby('Cluster').mean()
print(cluster_analysis)
```



Cluster	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
-1	0.550000	39.716667	64.216667	33.083333
0	0.588235	24.764706	25.588235	74.764706
1	0.573171	43.024390	55.109756	49.304878
2	0.548387	32.709677	79.774194	82.483871
3	0.500000	42.400000	83.200000	18.400000

```
# Cluster Plot
```

```
plt.figure(figsize=(8, 6))
plt.scatter('Annual Income (k$)', 'Spending Score (1-100)', data=features_DBSCAN, c='Cluster', cmap='viridis', s=50, alpha=0.5)
plt.title('Clusters de Clientes - DBSCAN')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.colorbar(label='Cluster DBSCAN')
plt.show()
```



Agglomerative Clustering

```
# Finding the best Cluster Number

features_AC = df_segmentation.iloc[:, 1:] # Remove Id
def find_optimal_clusters(features_DBSCAN, max_clusters=10):
    silhouette_scores = []

    for n_clusters in range(2, max_clusters+1):
        model = AgglomerativeClustering(n_clusters=n_clusters)
        labels = model.fit_predict(features_AC)
        silhouette_avg = silhouette_score(features_AC, labels)
        silhouette_scores.append(silhouette_avg)


    # Plot the silhouette scores

    plt.plot(range(2, max_clusters+1), silhouette_scores)
    plt.xlabel('Number of Clusters')
    plt.ylabel('Silhouette Score')
    plt.title('Silhouette Score for Different Number of Clusters')
    plt.show()

    optimal_clusters = np.argmax(silhouette_scores) + 2 # +2 because the range starts at 2

    return optimal_clusters

optimal_clusters = find_optimal_clusters(features_DBSCAN, max_clusters=10)
print("Best Cluster Number:", optimal_clusters)
```



Best Cluster Number: 6

```
# Setting Agglomerative Clustering
```

```
model = AgglomerativeClustering(n_clusters=5) # I will use the cluster number 5 due to the silhouette score plot
Clusters_AC = model.fit_predict(features_kmeans)
features_AC.loc[:, 'Cluster'] = Clusters_AC
```

```
# Analyze cluster characteristics
```

```
cluster_analysis = features_AC.groupby('Cluster').mean()
print(cluster_analysis)
```



Cluster	Gender	Age	Annual Income (k\$)	Spending Score (1-100)
0	0.602410	42.156627	54.759036	49.771084
1	0.428571	41.685714	88.228571	17.285714
2	0.538462	32.692308	86.538462	82.128205
3	0.600000	24.850000	24.950000	81.000000
4	0.608696	45.217391	26.304348	20.913043

```
# Cluster Plot
```

```
plt.figure(figsize=(8, 6))
plt.scatter('Annual Income (k$)', 'Spending Score (1-100)', data=features_AC, c='Cluster', cmap='viridis', s=50, alpha=0.5)
plt.title('Clusters de Clientes - Agglomerative Clustering')
plt.xlabel('Annual Income (k$)')
plt.ylabel('Spending Score (1-100)')
plt.colorbar(label='Cluster Agglomerative Clustering')
plt.show()
```




Clusters de Clientes - Agglomerative Clustering



Model Evaluation



Setting Models

```
n_clusters = 5
```

```
models = {
    'KMeans': KMeans(n_clusters=n_clusters, random_state=42, n_init = 10),
    'DBSCAN': DBSCAN(eps=13, min_samples=10), # same parameters from previous model
    'AgglomerativeClustering': AgglomerativeClustering(n_clusters=n_clusters)
}
```



Clustering Evaluation with Silhouette Score and Davies Bouldin Index

```
silhouette_scores = {}
davies_bouldin_indices = {}
labels_dict = {}
features = df_segmentation.iloc[:, 1:]
```

```
for name, model in models.items():
    # Train
    labels = model.fit_predict(features)
    labels_dict[name] = labels

    # Performance
    silhouette_avg = silhouette_score(features, labels)
    db_index = davies_bouldin_score(features, labels)

    silhouette_scores[name] = silhouette_avg
    davies_bouldin_indices[name] = db_index

    print(f'{name} - Silhouette Score: {silhouette_avg}, Davies-Bouldin Index: {db_index}')
```



```
KMeans - Silhouette Score: 0.44424291275274114, Davies-Bouldin Index: 0.8196433344356222
DBSCAN - Silhouette Score: 0.18699358696080978, Davies-Bouldin Index: 2.28192907500783
AgglomerativeClustering - Silhouette Score: 0.43997527212476695, Davies-Bouldin Index: 0.8220436090843712
```

Plot

```
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
```

```
for ax, (name, labels) in zip(axes, labels_dict.items()):
    ax.scatter('Annual Income (k$)', 'Spending Score (1-100)', data=features, c=labels, cmap='viridis', marker='o')
    ax.set_title(f'{name} Clustering')
    ax.set_xlabel('Annual Income (k$)')
    ax.set_ylabel('Spending Score (1-100)')
```

```
plt.show()
```

